

# Spire

by example

Code: <http://github.com/non/spire>

Learn: <http://typelevel.org>

Tom Switzer  
@tixxit





sign in new guest? my account REDcard

all search



women men baby kids home furniture electronics entertainment toys beauty health halloween deals more

REDcard



SAVE 5% + GET FREE SHIPPING.  
TODAY & EVERYDAY - Apply today

find a store

Weekly Ad

GiftCards

registries

TargetLists



Target > baby > car seats > harness booster car seats

sale

## Graco Nautilus Car Seat - Luxe



202 reviews

[write a review](#)

**\$134.99**

Reg: \$149.99 - Save \$15.00

• Free shipping when you spend \$50

quantity: \*

1



add to cart

× Not sold in stores

[add to registry](#)

[add to list](#)



## popular picks for baby.



### Up&Up™ Baby Diapers - Bulk ...

\$25.99 - \$26.99

★★★★★ (358)

• Free shipping when you spend \$50



### Toddler Girl's Circo® Das...

\$24.99

# The Problem

Some things depend on a set of **weights**.

Weights were chosen “**heuristically**.”


Can we optimize the weights?

# Tools


Split traffic to run A/B tests on weights.

Calculate metrics like RPS.


Is this enough?




# Gradient Descent



# Gradient Descent




# Gradient Descent





# Gradient Descent



# Requirements for Optimizing F

Need a starting position:  $w_0$

Need a “step” size (learning rate):  $\alpha$

Need to calculate the gradient:  $\nabla F$

$$w_{n+1} = w_n - \alpha \nabla F(w_n)$$

# In a Language We Understand

```
trait GradientDescent {  
  type V = Array[Double]  
  type A = Double  
  
  def grad(i: Int, w: V): V  
  
  def nextWeight(i: Int, w: V, a: A): V =  
    w - (a *: grad(i, w))  
}
```

# In a Language We Understand

```
trait GradientDescent {
```

```
  type V = Array[Double]
```

← the vector type

```
  type A = Double
```

← the scalar type

```
  def grad(i: Int, w: V): V
```

```
  def nextWeight(i: Int, w: V, a: A): V =
```

```
    w - (a *: grad(i, w))
```

←  $w_{n+1} = w_n - \alpha \nabla F(w_n)$

```
}
```

# What Spire Gave Us

```
trait GradientDescent {  
  type V = Array[Double]  
  type A = Double
```

```
  def grad(i: Int, w: V): V
```

Multiply vector by a scalar

```
  def nextWeight(i: Int, w: V, a: A): V =  
    w - (a *: grad(i, w))
```

```
}
```

# What Spire Gave Us

```
trait GradientDescent {  
  type V = Array[Double]  
  type A = Double
```

```
  def grad(i: Int, w: V): V
```

**Subtract 2 vectors**

```
  def nextWeight(i: Int, w: V, a: A): V =  
    w - (a *: grad(i, w))
```

```
}
```

# Why Does It Work?

`Double` is a “field”

`Array[Double]` is a “vector space”

We can add/subtract vectors.

And multiply them by scalars.





# Algebras Used In This Talk

`Field[A]`

`VectorSpace[V, A]`

`InnerProductSpace[V, A]`

`Order[A]`

# Field[A]

Has +, -, \*, and /.

Instances for Float, Double, BigDecimal, Rational, Complex, etc.

# Example

```
import spire.algebra.Field
```

```
import spire.implicits._
```

```
def mean[A: Field](xs: A*) =  
  xs.foldLeft(Field[A].zero)(_ + _) / xs.size
```

```
mean(1F, 2F, 3F)
```

```
mean(1D, 2D, 3D)
```

```
mean(r"1/2", r"2/3", r"4/5")
```

# Example

```
import spire.algebra.Field
import spire.implicit._
```

Provides instances  
and syntax



```
def mean[A: Field](xs: A*) =  
  xs.foldLeft(Field[A].zero) (_ + _) / xs.size
```

```
mean(1F, 2F, 3F)
```

```
mean(1D, 2D, 3D)
```

```
mean(r"1/2", r"2/3", r"4/5")
```

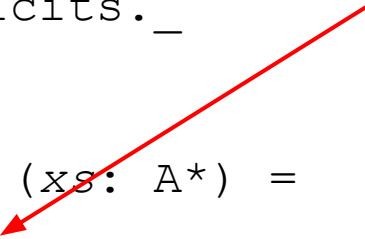
# Example

```
import spire.algebra.Field
import spire.implicits._
```

Short-hand for:

```
implicitly[Field[A]].
zero
```

```
def mean[A: Field](xs: A*) =  
  xs.foldLeft(Field[A].zero) (_ + _) / xs.size
```



```
mean(1F, 2F, 3F)
```

```
mean(1D, 2D, 3D)
```

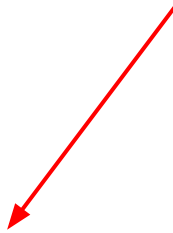
```
mean(r"1/2", r"2/3", r"4/5")
```

# Example

```
import spire.algebra.Field
import spire.implicits._
```

Division is exact

```
def mean[A: Field](xs: A*) =  
  xs.foldLeft(Field[A].zero) (_ + _) / xs.size
```



```
mean(1F, 2F, 3F)
```

```
mean(1D, 2D, 3D)
```

```
mean(r"1/2", r"2/3", r"4/5")
```

# Example

```
import spire.algebra.Field
import spire.implicits._
```

```
def mean[A: Field](xs: A*) =
  xs.foldLeft(Field[A].zero) (_ + _) / xs.size
```

```
mean(1F, 2F, 3F)
```

```
mean(1D, 2D, 3D)
```

```
mean(r"1/2", r"2/3", r"4/5")
```

Works on Float,  
Double, Rational,  
etc.

# **VectorSpace[V,A]**

Operations on vectors  $V$  with scalar  $A$ .  
The scalar  $A$  must be a Field.

Instances for Array, Vector, List, etc.



# InnerProductSpace[V,A]

Like VectorSpace, but adds an inner product.

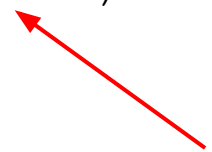
Vector(1D, 2D) dot Vector(3D, 4D)



# Order[A]

Provides a total ordering on A.

```
def sort[A: Order](xss: List[A]) = xss match {  
  case x :: xs =>  
    val (ls, rs) = xs.partition(_ <= x)  
    sort(ls) ++ (c :: rs)  
  case Nil => Nil  
}
```



Also has >, ==, cmp, min, etc.



**END  
DETOUR**

# Back to the Example

```
trait GradientDescent {
```

```
  type V = Array[Double]
```

```
  type A = Double
```



Do we really care what V and A are?

```
  def grad(i: Int, w: V): V
```

```
  def nextWeight(i: Int, w: V, a: A): V =
```

```
    w - (a *: grad(i, w))
```

```
}
```

# Making it Generic

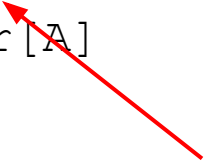
```
trait GradientDescent[V, A] {  
  implicit def V: VectorSpace[V, A]  
  implicit def A: Field[A]  
  
  def grad(i: Int, w: V): V  
  
  def nextWeight(i: Int, w: V, a: A): V =  
    w - (a *: grad(i, w))  
}
```

# The Optimization Loop

```
trait GradientDescent[V, A] {  
  ...  
  def optimize(a: A): V = ...  
  ...  
}
```

# The Optimization Loop

```
trait GradientDescent[V, A] {  
  implicit def V: InnerProductSpace[V, A]  
  implicit def order: Order[A]  
  ...  
  def optimize(a: A): V = ...  
  ...  
}
```

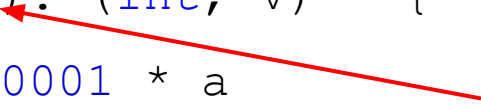


We'll need the dot product.

# The Optimization Loop

```
def optimize(a: A): (Int, V) = {  
    val eps = 0.0000001 * a  
    def loop(w0: V, i: Int): (Int, V) = {  
        val w1 = nextWeight(i, w0, a)  
        val dw = (w0 - w1)  
        if ((dw dot dw) > eps) loop(w1, i + 1) else (i, w1)  
    }  
    loop(V.zero, 0)  
}
```

Step/learning rate

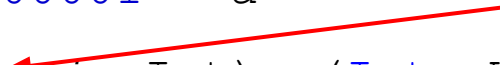




# The Optimization Loop

```
def optimize(a: A): (Int, V) = {  
    val eps = 0.0000001 * a  
    def loop(w0: V, i: Int): (Int, V) = {  
        val w1 = nextWeight(i, w0, a)  
        val dw = (w0 - w1)  
        if ((dw dot dw) > eps) loop(w1, i + 1) else (i, w1)  
    }  
    loop(V.zero, 0)  
}
```


Initial weight



# The Optimization Loop

```
def optimize(a: A): (Int, V) = {  
    val eps = 0.0000001 * a  
    def loop(w0: V, i: Int): (Int, V) = {  
        val w1 = nextWeight(i, w0, a)  
        val dw = (w0 - w1)  
        if ((dw dot dw) > eps) loop(w1, i + 1) else (i, w1)  
    }  
    loop(V.zero, 0)  
}
```

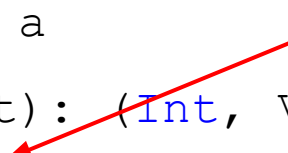
Current iteration



# The Optimization Loop

```
def optimize(a: A): (Int, V) = {  
    val eps = 0.0000001 * a  
    def loop(w0: V, i: Int): (Int, V) = {  
        val w1 = nextWeight(i, w0, a)  
        val dw = (w0 - w1)  
        if ((dw dot dw) > eps) loop(w1, i + 1) else (i, w1)  
    }  
    loop(V.zero, 0)  
}
```

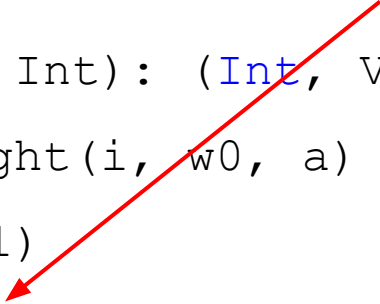
Perform 1 step



# The Optimization Loop

```
def optimize(a: A): (Int, V) = {  
    val eps = 0.0000001 * a  
    def loop(w0: V, i: Int): (Int, V) = {  
        val w1 = nextWeight(i, w0, a)  
        val dw = (w0 - w1)  
        if ((dw dot dw) > eps) loop(w1, i + 1) else (i, w1)  
    }  
    loop(V.zero, 0)  
}
```

Check if we've "converged"



# Testing it Out: Linear Least Squares

Problem: Fit a line to a set of points.

- Line:  $y = mx + b$
- Points:  $[(x_0, y_0), (x_1, y_1), \dots, (x_n, y_n)]$

Minimize:  $\sum (mx_i + b - y_i)^2$

# Gradient Descent Requirements

- ✓ Need a starting position.
- ✓ Need a “step” size (learning rate).
- ✓ Need to calculate the **gradient**.

# The Derivative: 2 Dimensions

$$F(w) = \sum (mx_i + b - y_i)^2$$

$$f_i(m, b) = (mx_i + b - y_i)^2$$

$$f'_i(m, b) = [2x_i(mx_i + b - y_i), 2(mx_i + b - y_i)]$$

$$\nabla F(w) = \sum [2x_i(mx_i + b - y_i), 2(mx_i + b - y_i)]$$

# The Derivative: N Dimensions

$$p = [x_0, x_1, \dots, x_n, y]$$

$$\beta = [x_0, x_1, \dots, x_n, 1]$$

$$f_p(w) = (\beta \cdot w - y)^2$$

$$\begin{aligned} f'_p(w) &= [2x_0(\beta \cdot w - y), 2x_1(\beta \cdot w - y), \dots, 2(\beta \cdot w - y)] \\ &= \mathbf{2(\beta \cdot w - y)\beta} \end{aligned}$$

$$\nabla F(w) = \sum 2(\beta_p \cdot w - y_p)\beta_p$$



# In Spire

```
class LeastSquares[V, A](betas: Iterable[V], ys: Iterable[A])
  extends GradientDescent[V, A] {


  val grads: Iterable[V => V] =
    (betas zip ys) map { case (b, y) =>
      { (w: V) => (2 * ((b dot w) - y)) *: b }
    }

  def grad(i: Int, v: V): V =
    grads.foldLeft(V.zero)(_ + _(v)) :/ grads.size
}
```

# In Spire

```
class LeastSquares[V, A](betas: Iterable[V], ys: Iterable[A])  
  extends GradientDescent[V, A] {
```

```
  val grads: Iterable[V => V] =  
    (betas zip ys) map { case (b, y) =>  
      { (w: V) => (2 * ((b dot w) - y)) *: b }  
    }
```


$$2(\beta \cdot w - y)\beta$$

```
  def grad(i: Int, v: V): V =  
    grads.foldLeft(V.zero)(_ + _(v)) :/ grads.size  
}
```

# In Spire

```
class LeastSquares[V, A](betas: Iterable[V], ys: Iterable[A])  
  extends GradientDescent[V, A] {
```

```
  val grads: Iterable[V => V] =  
    (betas zip ys) map { case (b, y) =>  
      { (w: V) => (2 * ((b dot w) - y)) *: b }  
    }  
}
```

Scalar division



```
def grad(i: Int, v: V): V =  
  grads.foldLeft(V.zero) (_ + _(v)) :/ grads.size  
}
```

# Testing it Out

```
import spire.random.Dist
```

```
def linear[A: ...](coeffs: Array[A]): Dist[Array[A]] = ...
```

```
val points = linear(Array(0.7, 2.3)).sample(40)
```

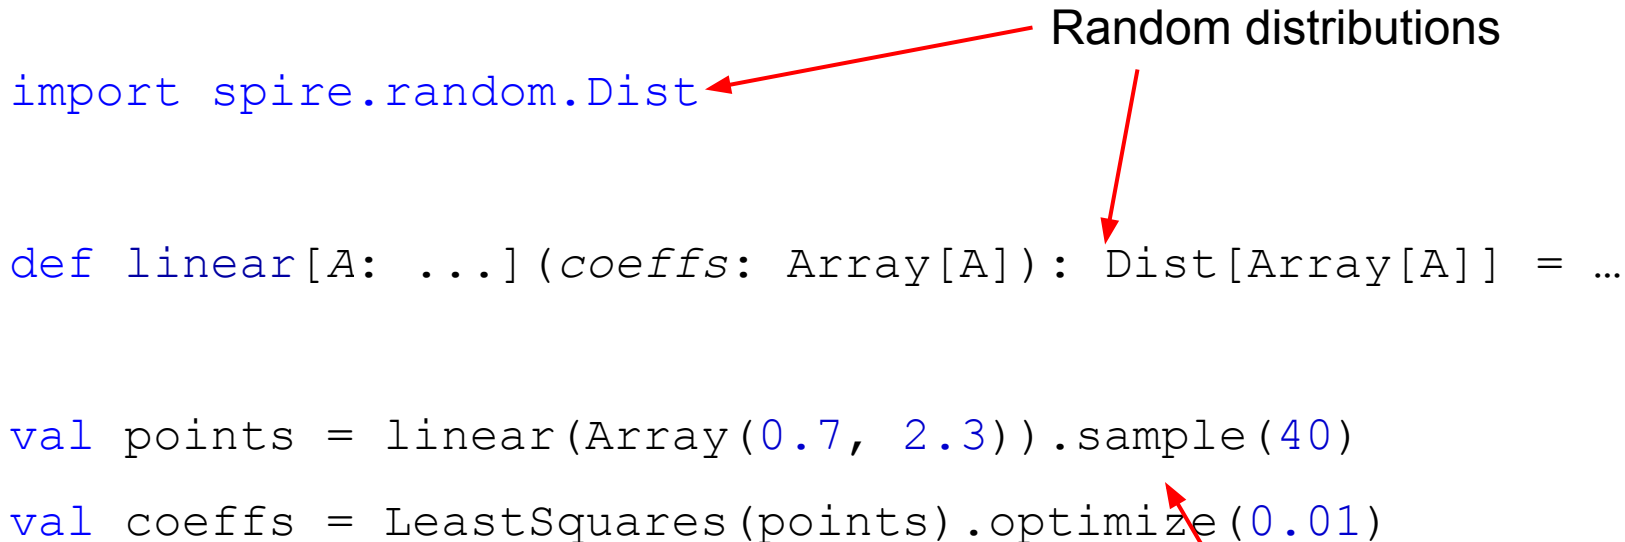
```
val coeffs = LeastSquares(points).optimize(0.01)
```

# Testing it Out

```
import spire.random.Dist

def linear[A: ...](coeffs: Array[A]): Dist[Array[A]] = ...

val points = linear(Array(0.7, 2.3)).sample(40)
val coeffs = LeastSquares(points).optimize(0.01)
```



Create a sample with 40 points

# Testing it Out

```
import spire.random.Dist
```


```
def linear[A: ...](coeffs: Array[A]): Dist[Array[A]] = ...
```

```
val points = linear(Array(0.7, 2.3)).sample(40)
```

```
val coeffs = LeastSquares(points).optimize(0.01)
```

Optimize using gradient descent!





$$f(x) = 0.8x + 1.76$$

# Back to Our Actual Problem

Optimize “weights” in personalization product.

Can we use gradient descent?



# Requirements

- ✓ Need a starting position.
- ✓ Need a “step” size (learning rate).
- ✗ Need to calculate the **gradient**.

# Requirements

- ✓ Need a starting position.
- ✓ Need a “step” size (learning rate).
- ✗ Need to calculate the **gradient**.

**Unknown  $F$ , unknown  $\nabla F$**

# Approximate the Gradient

Sample around  $w$ .

Fit hyperplane (line) to samples.

The coefficients approximate the gradient.

# Approximate the Gradient

Sample around  $w$ .

Fit hyperplane (line) to samples.

The coefficients approximate the gradient.

# Approximate the Gradient

```
abstract class SampledGradientDescent[A: ClassTag: Uniform]  
  (f: (Int, Array[A]) => A, llsLearningRate: A)  
  extends GradientDescent[Array[A], A] {  
    def timespan(i: Int): Range  
    def sample(v: Array[A], d: A): Array[Array[A]] = ...  
    def grad(i: Int, v: Array[A]): Array[A] = ...  
  }
```

# Approximate the Gradient

...

```
def sample(v: Array[A], d: A): Array[Array[A]] = {  
  import Dist._  
  val noise = array(v.length) (uniform[A] (-d, d))  
  (constant(v) + noise)  
    .sample[Vector] (v.length + 1)  
    .toArray  
}
```

...

# Approximate the Gradient

...

```
def sample(v: Array[A], d: A): Array[Array[A]] = {  
  import Dist._  
  val noise = array(v.length) (uniform[A] (-d, d))  
  (constant(v) + noise)  
    .sample[Vector](v.length + 1)  
    .toArray  
}
```

Methods from Dist.\_

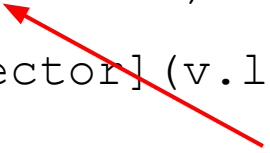


...

# Approximate the Gradient

...

```
def sample(v: Array[A], d: A): Array[Array[A]] = {  
  import Dist._  
  val noise = array(v.length) (uniform[A] (-d, d))  
  (constant(v) + noise)  
    .sample[Vector] (v.length + 1)  
    .toArray  
}
```



Adding “distributions” together



# Approximate the Gradient

...

```
def grad(i: Int, v: Array[A]): Array[A] = {  
    val vs = sample(v, llsLearningRate * 4)  
    val g = LeastSquares(vs, vs map { v =>  
        timespan(i).map(f(_, v)).qmean  
    }).optimize(llsLearningRate)._2.init  
}
```

...

# Approximate the Gradient

...

```
def grad(i: Int, v: Array[A]): Array[A] = {  
  val vs = sample(v, llsLearningRate * 4)  
  val g = LeastSquares(vs, vs map { v =>  
    timespan(i).map(f(_, v)).qmean  
  }).optimize(llsLearningRate)._2.init  
}
```

...

Use our linear least squares  
implementation to fit hyperplane!

# Approximate the Gradient

...

```
def grad(i: Int, v: Array[A]): Array[A] = {  
  val vs = sample(v, llsLearningRate * 4)  
  val g = LeastSquares(vs, vs map { v =>  
    timespan(i).map(f(_, v)).qmean  
  }).optimize(llsLearningRate)._2.init  
}
```

...


Average over a span of time



# Approximate the Gradient

...

```
def grad(i: Int, v: Array[A]): Array[A] = {  
  val vs = sample(v, llsLearningRate * 4)  
  val g = LeastSquares(vs, vs map { v =>  
    timespan(i).map(f(_, v)).qmean  
  }).optimize(llsLearningRate)._2.init  
}
```



...

Drop the value axis intercept

# **So... What's the “timespan” Thing?**

People shop differently on different days.

Time is an implicit parameter.

For testing, make it explicit.

# A New Challenger Arrives

Evaluating  $F$  at  $w$ :

- Split traffic: some see weights  $w$
- Wait  $X$  hours/days/weeks
- Calculate our metric

# A New Challenger Arrives

Evaluating  $F$  at  $w$ :

- Split traffic: some see weights  $w$
- Wait  $X$  hours/days/weeks
- Calculate our metric



**Convergence may take a while!**

# A New Challenger Arrives

Evaluating  $F$  at  $w$ :

- Split traffic: some see weights  $w$
- Wait  $X$  hours/days/weeks
- Calculate our metric

How does  $X$  affect the ***time*** to convergence?



# Test It Out!

Our Assumption:

- $F$  varies periodically over a week

# Test It Out!

Our Assumption:

- F varies periodically over a week

Strategies:

- Fix X to 1 or 7
- Linear ramp-up from 1-7
- Run it at 1 at first, then jump to 7

# Strategy 1: Fix X

```
def period: Int
```

```
def timespan(i: Int): Range =  
    (i * period) until ((i + 1) * period)
```

# Strategy 2: Ramp-up X

```
def iters: Int
def max: Int
def timespan(i: Int): Range = {
    val p = spire.math.min(max, (i / iters) + 1)
    val j = i - (p - 1) * iters
    val start = p * j + iters * (p * (p - 1)) / 2
    val end = start + p
    start until end
}
```

# Strategy 2: Ramp-up X

```
def iters: Int
```

```
def max: Int
```

```
def timespan(i: Int): Range = {
```

```
    val p = spire.math.min(max, (i / iters) + 1)
```

```
    val j = i - (p - 1) * iters
```

```
    val start = p * j + iters * (p * (p - 1)) / 2
```

```
    val end = start + p
```

```
    start until end
```

```
}
```

Note: not scala.math!



# Strategy 3: Jump Up

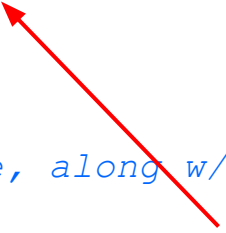
```
def fst: Int
def snd: Int
def jump: Int
def timespan(i: Int): Range = {
  if (i < jump) {
    (i * fst) until ((i + 1) * fst)
  } else {
    val start = jump * fst + (i - jump) * snd
    start until (start + snd)
  }
}
```

# The Data: Creating a Cyclic F

```
def cyclic[A: Field: Order](p: Int,  
    ps: (Int, Array[A] => A)*): (Int, Array[A]) => A = {  
    val (indices, fs) = ps.toArray.qsortedBy(_._1).unzip  
  
    def eval(t: Int, v: Array[A]): A = {  
        val i = indices.qsearch(t) // Note, along w/ qselect and qsort.  
        ...  
    }  
  
    { (t, v) => eval(t % period, v) }  
}
```

# The Data: Creating a Cyclic F

```
def cyclic[A: Field: Order](p: Int,  
    ps: (Int, Array[A] => A)*): (Int, Array[A]) => A = {  
    val (indices, fs) = ps.toArray.qsortedBy(_._1).unzip  
  
    def eval(t: Int, v: Array[A]): A = {  
        val i = indices.qsearch(t) // Note, along w/ qselect and qsort.  
        ...  
    }  
  
    { (t, v) => eval(t % period, v) }  
}
```




Sorting!



# The Data: Creating a Cyclic F

```
def cyclic[A: Field: Order](p: Int,  
    ps: (Int, Array[A] => A)*): (Int, Array[A]) => A = {  
    val (indices, fs) = ps.toArray.qsortedBy(_._1).unzip  
  
    def eval(t: Int, v: Array[A]): A = {  
        val i = indices.qsearch(t) // Note, along w/ qselect and qsort.  
        ...  
    }  
  
    { (t, v) => eval(t % period, v) }  
}
```



Binary search!

# Other Methods on Seq and Array

Sorting (`xs.qsort`)

Searching (`xs.qsearch(x)`)

- Never implement binary search again

Selection (`xs.qselect(k)`)

- Linear time selection

# A Week in Code

```
def quad[A: Field](a: Double, h: Double, k: Double): Array[A] => A =  
  { v => a * ((v(0) - h) ** 2) + k }
```

```
def weekly[A: Field: Order]: (Int, Array[A]) => A =  
  cyclic(7, 0 -> quad(2 , 2 , 5),  
          1 -> quad(1.5, 2.5, 3),  
          3 -> quad(1 , 3 , 1),  
          5 -> quad(1.5, 2.5, 3),  
          6 -> quad(2 , 2 , 5))
```

# A Week in Code

```
def quad[A: Field](a: Double, h: Double, k: Double): Array[A] => A =  
  { v => a * ((v(0) - h) ** 2) + k }
```

```
def weekly[A: Field: Order]: (Int, Array[A]) => A =  
  cyclic(7, 0 -> quad(2 , 2 , 5),  
        1 -> quad(1.5, 2.5, 3),  
        3 -> quad(1 , 3 , 1),  
        5 -> quad(1.5, 2.5, 3),  
        6 -> quad(2 , 2 , 5))
```

The weekend



# A Week in Code

```
def quad[A: Field](a: Double, h: Double, k: Double): Array[A] => A =  
  { v => a * ((v(0) - h) ** 2) + k }
```

```
def weekly[A: Field: Order]: (Int, Array[A]) => A =
```

```
  cyclic(7, 0 -> quad(2 , 2 , 5),  
          1 -> quad(1.5, 2.5, 3),  
          3 -> quad(1 , 3 , 1),  
          5 -> quad(1.5, 2.5, 3),  
          6 -> quad(2 , 2 , 5))
```

The weekend

Minimum over the  
week is ~2.3

# What We Want

Minimize # of days to convergence.

Minimize error of result.

# What We Want

Minimize # of days to convergence.

Minimize error of result.

These goals may compete with each other!

# Results: No Free Lunch

Strategy	# of iterations	# of days	Converged To	RMSE
<b>Fixed: 7 days</b>	475	3325	2.245	0.73
<b>Fixed: 1 day</b>	1189	1189	2.21	<b>1.23</b>
<b>Ramp-up: 1-7</b>	727	2659	2.23	0.99
<b>Jump: 200</b>	569	2789	2.27	0.71
<b>Jump: 300</b>	559	2119	2.26	0.88
<b>Jump: 400</b>	727	2659	2.23	0.99



# Results: No Free Lunch

Strategy	# of iterations	# of days	Converged To	RMSE
<b>Fixed: 7 days</b>	475	3325	2.245	0.73
<b>Fixed: 1 day</b>	1189	1189	2.21	1.23
<b>Ramp-up: 1-7</b>	727	2659	2.23	0.99
<b>Jump: 200</b>	569	2789	2.27	0.71
<b>Jump: 300</b>	559	2119	2.26	0.88
<b>Jump: 400</b>	727	2659	2.23	0.99

**Questions?**